

50 CareerFlow Interview Q&A

Conversational answers for AI Engineer Round 2 (onsite / deep dive)

CareersLow x CareerFlow AI Engineer - Round 2 onsite prep. Conversational answers with trade-offs, eval, latency/cost, and reroutes.

Evaluation & Quality

1. How do you evaluate an LLM feature before shipping it?

In Ca
retriev
sugge
requir
full pip
end-to
upfront
wire L
score

2. What's in your golden dataset?

It live
rough
shoul
layer
regres
labelin
6.2 ve

3. What metrics do you use for retrieval?

Mainly
find a
were a
good
sixty-t
labeled
label r

4. What metrics do you use for gap scoring?

Bucke
wheth
behav
6.5; th
MAE v

5. Why evaluate retrieval separately if you already have end-to-end eval?

Because
looks
and th
any e
of only

6. How do you know your eval dataset isn't garbage?

First,
has n
not ca
are no

7. What's your pass/fail bar?

For re
accur
lever
report
Hit at

8. How would you detect hallucinations in production?

In my
eviden
where
add s
trade-
scoring

9. How do you monitor model drift?

I do n
datas
JSON
golden
locally

10. How do you test mock interview quality?

Voice
a sess
- no
that w

11. How do you regression-test prompt changes?

Any p
saved
non-d
contro
- so y

12. What would you add to eval next?

Suggest
judgment
would
value.

13. How do you handle flaky LLM eval runs?

Score
rate li
feedb

14. How is your eval approach relevant to CareerFlow's JD?

The J
built c
transf
LangS
direct

15. Walk me through one eval failure you found and fixed.

Duties
Skills
are va
fix is t
the la

Latency & Cost

16. Where is latency in your pipeline?

Extrac
LangC
analys
stream
pair.

17. How does Redis reduce cost?

I cach
by res
The J
withou
would

18. Why skip Qdrant re-indexing?

On re
or fifty
- that

19. Why parallel gap scoring with a semaphore?

20. Why decomposed scorer and writer instead of one LLM call?

21. How would you reduce token cost further?

22. How do you optimize mock interview latency?

23. Why Gemini instead of OpenAI - cost implications?

24. What's your caching invalidation strategy?

25. How would you design this for 1.2M users on Cloud Run?

RAG & Retrieval

26. Why only index the resume, not the JD?

27. Why hybrid dense plus BM25?

relying
keywo
scale:

This c
test a
the ex
says :
miss l
BM25
two re

28. Why chunk at list-item level?

When
section
docum
retriev
on pa
tuning

29. Why section filters for skills versus duties?

I add
summ
eviden
exper
trade-
precis

30. Why a score threshold on dense retrieval?

So wh
becau
loosel
eviden
phras

31. Why delete before upsert in Qdrant?

Imagi
new c
So be
delete

32. Why not fine-tune instead of RAG?

Hones
differe
why s
perso
cost a

33. How would you use spaCy here?

I did not
layout
already
and so

LLM Pipeline & LangGraph

34. Why LangGraph for this pipeline?

I chose
parse
index
index
match

35. Why structured output with Pydantic?

Produ
GapA
to val
repair

36. Why temperature zero for scoring and zero point four for writing?

Scoring
same
read
separ

37. Why sixty-five, fifteen, twenty blend for overall fit?

Requi
skills
up wi
wheth

38. Why cap importance weights at nine?

When
does
HAVE

39. Why batch bullet generation?

After
sugge
The tr

40. How would you add self-correction?

I did not
If Pyd
groun

Voice & Real-time

41. Why WebSocket instead of REST for mock interview?

A mock
transcription
user
operation

42. What is VAD and why did you need it?

VAD is
speech
timer
would
trade-

43. What is barge-in and how did you handle it?

Barge
browser
meta
the AI

44. Why Sarvam not Deepgram or ElevenLabs?

Career
the ha
uttera
plug in

45. How does voice connect to gap analysis?

Modu
scores
the vo
record
fully fr

Production & System Design

46. What happens when Redis is down?

I tried
metho
and m

47. How do you handle PDF extraction failures?

First s
with a
resum

48. How would you ship a prompt change to 1.2M users?

Run m
traffic
move
system

49. Design CareerFlow's multi-agent system.

I woul
Firest
bullet
their c
end to

50. Why should we hire you for this role?

Caree
optimi
the sa
separ
issues
prod,
the pr